



Module 7: Advanced SQL

Edited by Radhika Sukapuram

Database System Concepts, 7th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Outline

- Accessing SQL From a Programming Language
- Functions and Procedures
- Triggers



Accessing SQL from a Programming Language

A database programmer must have access to a general-purpose programming language. Why ?

- Not all queries can be expressed in SQL
- Non-declarative actions
 - printing a report
 - interacting with a user
 - sending the results of a query to a graphical user interface

cannot be done from SQL



Accessing SQL from a Programming Language (Cont.)

There are two approaches to accessing SQL from a general-purpose programming language

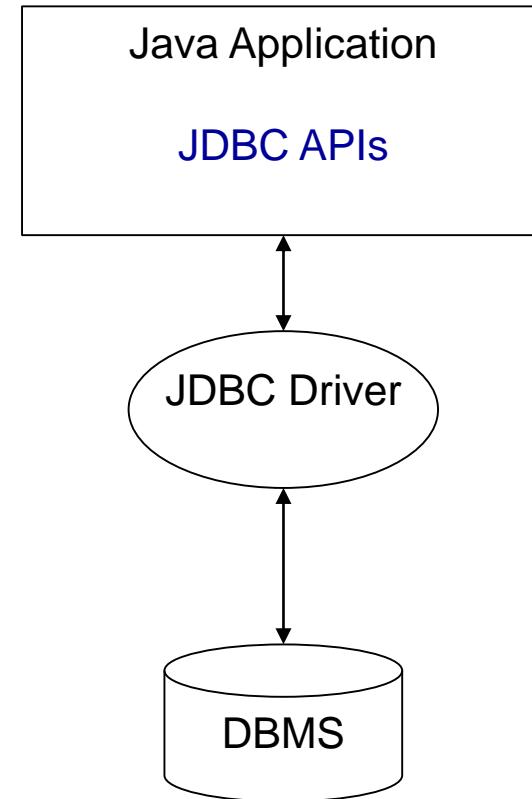
- Dynamic SQL: A general-purpose program
 - can connect to and communicate with a database server using a collection of functions/methods
 - Construct an SQL query as a character string and submit
 - Receive results into variables a tuple at a time
 - Two standards specifying APIs
 - ▶ Java Database Connectivity
 - ▶ Open Database Connectivity
- Embedded SQL
 - SQL statements are embedded
 - translated at compile time into function calls
 - At runtime, these function calls connect to the database using an API that provides dynamic SQL facilities.



Java Database Connectivity: JDBC



JDBC





JDBC

- JDBC is a Java API for communicating with database systems supporting SQL.
- JDBC supports
 - querying and updating data
 - retrieving query results
 - metadata retrieval
 - ▶ querying about relations present in the database and the names and types of relation attributes.
- Model for communicating with the database:
 - Open a connection
 - Create a “statement” object
 - Execute queries using the Statement object to send queries and fetch results
 - Exception mechanism to handle errors



JDBC Code

```
public static void JDBCexample(String userid, String passwd)
{
    try (Connection conn = DriverManager.getConnection(
        "jdbc:oracle:thin:@db.yale.edu:2000:univdb", userid, passwd);
        Statement stmt = conn.createStatement()); //To send SQL statements
    )
    {
        /*... Do Actual Work .... */
    }
    catch (SQLException sqle) {
        System.out.println("SQLException : " + sqle);
    }
}
```

NOTE: Above syntax works with Java 7, and JDBC 4 onwards.

Resources opened in “try (...)” syntax are automatically closed at the end of the try block



JDBC Code for Older Versions of Java/JDBC

```
public static void JDBCexample(String userid, String passwd)
{
    try {
        Class.forName ("oracle.jdbc.driver.OracleDriver");
        Connection conn = DriverManager.getConnection(
            "jdbc:oracle:thin:@db.yale.edu:2000:univdb", userid, passwd);
        Statement stmt = conn.createStatement();
        ... Do Actual Work ....
        stmt.close();
        conn.close();
    }
    catch (SQLException sqle) {
        System.out.println("SQLException : " + sqle);
    }
}
```

NOTE: Class.forName is not required from JDBC 4 onwards. The try with resources syntax in previous slide is preferred for Java 7 onwards.



JDBC Code (Cont.)

- Update to database

```
try {  
    stmt.executeUpdate(  
        "insert into instructor values(' 77987' , ' Kim' , ' Physics' ,  
        98000)");  
} catch (SQLException sqle)  
{  
    System.out.println("Could not insert tuple. " + sqle);  
}
```

- Execute query and fetch and print results

```
ResultSet rset = stmt.executeQuery(  
    "select dept_name, avg (salary)  
     from instructor  
     group by dept_name");  
  
while (rset.next()) {  
    System.out.println(rset.getString("dept_name") + " " +  
                      rset.getFloat(2)); //Attribute Position  
}
```



JDBC SUBSECTIONS

- Connecting to the Database
- Shipping SQL Statements to the Database System
- Exceptions and Resource Management
- Retrieving the Result of a Query
- Prepared Statements
- Callable Statements
- Metadata Features
- Other Features



JDBC Code Details

- Getting result fields:
 - **rs.getString(“dept_name”) and rs.getString(1) are equivalent if dept_name is the first argument of select result.**
- Dealing with Null values

```
int a = rs.getInt("a");
if (rs.wasNull()) System.out.println("Got null value");
```



Prepared Statement

- ```
PreparedStatement pStmt = conn.prepareStatement(
 "insert into instructor values(?, ?, ?, ?)");

pStmt.setString(1, "88877");
pStmt.setString(2, "Perry");
pStmt.setString(3, "Finance");
pStmt.setInt(4, 125000);
pStmt.executeUpdate();

pStmt.setString(1, "88878");
pStmt.executeUpdate();
```
- WARNING: always use prepared statements when taking an input from the user and adding it to a query
  - NEVER create a query by concatenating strings
  - **stmt.executeUpdate( “insert into instructor values(’ ” + ID + “ ’, ’ ” + name + “ ’, ’ ” + dept name + “ ’, ’ ”+ balance + ” )” );**
  - What if name is “D’ Souza”?



# SQL Injection

- Suppose query is constructed using
  - "select \* from instructor where name = ' " + name + " '"
- Suppose the user, instead of entering a name, enters:
  - X' or 'Y' = 'Y
- then the resulting statement becomes:
  - "select \* from instructor where name = ' " + "X' or 'Y' = 'Y" + " "
  - which is:
    - ▶ select \* from instructor where name = 'X' or 'Y' = 'Y'
  - User could have even used
    - ▶ X'; update instructor set salary = salary + 10000; --
- Prepared statement internally uses:  
"select \* from instructor where name = 'X\' or \'Y\' = \'Y'
  - **Always use prepared statements, with user inputs as parameters**



# Metadata Features

- ResultSet metadata
  - An object that can be used to get information about the types and properties of the columns in a ResultSet object.
- E.g. after executing query to get a ResultSet rs:
  - **ResultSet rs = stmt.executeQuery(...);**  
ResultSetMetaData rsmd = rs.getMetaData();  

```
for(int i = 1; i <= rsmd.getColumnCount(); i++) {
 System.out.println(rsmd.getColumnName(i));
 System.out.println(rsmd.getColumnTypeName(i));
}
```
- How is this useful?
  - Possible to execute query without knowing the schema of the result



# Metadata (Cont)

- Database metadata

- **Connection conn = DriverManager.getConnection(..);**

```
DatabaseMetaData dbmd = conn.getMetaData();
```

```
// Arguments to getColumns: Catalog, Schema-pattern, Table-pattern,
// and Column-Pattern
```

```
// Returns: One row for each column; row has a number of attributes
// such as COLUMN_NAME, TYPE_NAME
```

```
// The value null indicates all Catalogs/Schemas.
```

```
// The value "" indicates current catalog/schema. It is possible to set a
// schema
```

```
// The value "%" has the same meaning as SQL like clause
```

```
ResultSet rs = dbmd.getColumns(null, "univdb", "department", "%");
```

```
while(rs.next()) {
```

```
 System.out.println(rs.getString("COLUMN_NAME"),
 rs.getString("TYPE_NAME"));
```

```
}
```



# Metadata (Cont)

- Database metadata

- DatabaseMetaData dbmd = conn.getMetaData();  
    // Arguments to getTables: Catalog, Schema-pattern, Table-pattern,  
    // and Table-Type  
    // Returns: One row for each table; row has a number of attributes  
    // such as TABLE\_NAME, TABLE\_CAT, TABLE\_TYPE, ..  
    // The value null indicates all Catalogs/Schemas.  
    // The value "" indicates current catalog/schema  
    // The value "%" has the same meaning as SQL **like** clause  
    // The last attribute is an array of types of tables to return.  
    //   TABLE means only regular tables (not VIEWS)

```
ResultSet rs = dbmd.getTables ("", "", "%", new String[] {"TABLE"});
while(rs.next()) {
 System.out.println(rs.getString("TABLE_NAME"));
}
```



# Finding Primary Keys

```
■ DatabaseMetaData dmd = connection.getMetaData();

// Arguments below are: Catalog, Schema, and Table
// The value "" for Catalog/Schema indicates current catalog/schema
// The value null indicates all catalogs/schemas
// Retrieves a description of the given table's primary key columns.
//They are ordered by COLUMN_NAME
ResultSet rs = dmd.getPrimaryKeys("", "", tableName);

while(rs.next()){
 // KEY_SEQ indicates the position of the attribute in
 // the primary key, which is required if a primary key has multiple
 // attributes
 System.out.println(rs.getString("KEY_SEQ"),
 rs.getString("COLUMN_NAME"));
}
```



# Use of extracting metadata

- For tasks such as writing a database browser
- Making code for such tasks generic



# Transaction Control in JDBC

- By default, each SQL statement is treated as a separate transaction that is committed automatically
  - bad idea for transactions with multiple updates
- Can turn off automatic commit on a connection
  - `conn.setAutoCommit(false);`
- Transactions must then be committed or rolled back explicitly
  - `conn.commit();`    or
  - `conn.rollback();`
- `conn.setAutoCommit(true)` turns on automatic commit.



# Other JDBC Features

- Handling large object types
  - `getBlob()` and `getClob()` that are similar to the `getString()` method, but return objects of type Blob and Clob, respectively -- (ResultSet)
    - ▶ get data from these objects by `getBytes()` -- (Blob)
  - associate an open stream with Java Blob or Clob object to update large objects (PreparedStatement)
    - ▶ `blob.setBlob(int parameterIndex, InputStream inputStream)`.
    - ▶ `parameterIndex`: the position of the attribute to write into



# JDBC Resources

## ■ JDBC

- Oracle: <https://docs.oracle.com/javase/tutorial/jdbc/index.html>
- MySQL: <https://dev.mysql.com/doc/connector-j/8.0/en/>



# ODBC



# ODBC

- Open DataBase Connectivity (ODBC) standard
  - standard for application program to communicate with a database server.
  - application program interface (API) to
    - ▶ open a connection with a database,
    - ▶ send queries and updates,
    - ▶ get back results.
- Applications such as Microsoft spreadsheets, Microsoft Access etc. can use ODBC



# Embedded SQL



# Embedded SQL

- The SQL standard defines embeddings of SQL in a variety of programming languages such as C, C++, Java, Fortran, and PL/I
- A language in which SQL queries are embedded is referred to as a **host language**, and the SQL structures permitted in the host language comprise *embedded SQL*.
- **EXEC SQL** statement is used in the host language to identify embedded SQL request to the preprocessor

EXEC SQL <embedded SQL statement>;

Note: this varies by language:

- In Java embedding uses # sql { .... }; //Called SQLJ



# JDBC and Embedded SQL

Java program with JDBC

Interpreter

SQL statements  
are interpreted  
at run-time

Embedded SQL program

Special preprocessor

Replaces embedded SQL  
with host language  
declarations / procedural  
calls

Host language complier



SQL related errors may be caught at compile time



# Embedded SQL (Cont.)

- Before executing any SQL statements, the program must first connect to the database. This is done using:

EXEC-SQL **connect to** *server user user-name using password*;

Here, *server* identifies the server to which a connection is to be established.

- Variables of host language can be used in embedded SQL
  - The syntax for declaring the variables follows the usual host language syntax.

EXEC-SQL BEGIN DECLARE SECTION;

int *credit\_amount* ;

EXEC-SQL END DECLARE SECTION;

- To use in embedded SQL, prefix a colon

*:credit\_amount* = 100;



# Embedded SQL (Cont.)

- Incorporate the SQL error handling mechanism (SQL Communication Area)  
**EXEC SQL INCLUDE SQLCA;**
- Pre-defined variables are set to known values to indicate error



# Embedded SQL (Cont.)

- To write an embedded SQL query, we use the  
**declare c cursor for <SQL query>**  
statement. The variable *c* is used to identify the query
- Example:
  - From within a host language
    - find the ID and name of students who have completed more than the number of credits stored in variable *credit\_amount* in the host language
  - Specify the query in SQL as follows:

EXEC SQL

```
declare c cursor for
select ID, name
from student
where tot_cred > :credit_amount;
```



# Embedded SQL (Cont.)

- The query can be executed as:

**EXEC SQL open c ;**

- ▶ results are saved in a temporary relation.
- ▶ uses *credit\_amount* of the host language when executed.

- Place one tuple in the query result on host language variables:

**EXEC SQL fetch c into :si, :sn;**

- Repeated calls to fetch get successive tuples in the query result
- A variable SQLSTATE in SQLCA is set to ‘02000’ to indicate no more data is available



# Embedded SQL (Cont.)

- The **close** statement causes the database system to delete the temporary relation that holds the result of the query.

```
EXEC SQL close c ;
```

- The above details vary with the language.
  - Java embedding defines Java iterators to step through result tuples.



# Updates Through Embedded SQL

- Embedded SQL expressions for database modification (**update**, **insert**, and **delete**)
- Add 100 to the salary of every instructor in the Music department
- Can update tuples fetched by cursor by declaring that the cursor is for update

## EXEC SQL

```
declare c cursor for
select *
from instructor
where dept_name = 'Music'
for update;
```



# Updates Through Embedded SQL contd.

- We then iterate through the tuples by performing **fetch** operations on the cursor (as illustrated earlier), and after fetching each tuple we execute code to update :

```
EXEC SQL open c ;
do {
 EXEC SQL fetch c into :si, :sn, :sdept, :ssalary;
 if (ssalary < 10000){
 EXEC SQL
 update instructor
 set salary = salary + 100
 where current of c;}
} while !strcmp(SQLCA.SQLSTATE,'02000')) ;
```



# Functions and Procedures



# Functions and Stored Procedures

- Functions and procedures allow “business logic” to be stored in the database and executed from SQL statements
  - A rule on the minimum/maximum number of courses a student can take in a semester
  - Minimum number of courses an instructor must teach in a year
  - Maximum number of leaves an employee can take



# Functions and Procedures contd.

- Advantages
  - Minimizes data transfer between application and database
  - Utilizes power of DBMS server
  - Different users can reuse the stored procedure – reduces application logic
  - Needs change only in one place if business logic changes
- These can be defined either by
  - the procedural component of SQL or
  - by an external programming language - Java, C, or C++.
- The syntax we present here is defined by the SQL standard.
  - Most databases implement nonstandard versions
  - Oracle : PL/SQL, Microsoft SQL Server: TransactSQL , PostgreSQL(PL/pgSQL)
  - MySQL: <https://dev.mysql.com/doc/refman/8.0/en/create-procedure.html>



# Declaring SQL Functions

- Define a function:

- input - name of a department
  - output- number of instructors in that department.

```
create function dept_count (dept_name varchar(20))
 returns integer
begin
 declare d_count integer;
 select count (*) into d_count
 from instructor
 where instructor.dept_name = dept_name
 return d_count;
end
```



# Calling a function

- The function *dept\_count* can be used to
  - find the department names and budget of all departments with more than 12 instructors.

```
select dept_name, budget
from department
where dept_count(dept_name) > 12;
```
- A function can be used in an expression



# Table Functions

- **table functions** : functions that can return tables as results
- Example: Return all instructors in a given department

```
create function instructors_of(dept_name char(20))
```

```
 returns table (
```

```
 ID varchar(5),
 name varchar(20),
 dept_name varchar(20),
 salary numeric(8,2))
```

```
 return table
```

```
 (select ID, name, dept_name, salary
 from instructor
 where instructor.dept_name = instructors_of.dept_name)
```

- Usage

```
select *
from table (instructor_of(‘Music’));
```

Passed parameter

- Table functions can be thought of as **parameterized views**



# SQL Procedures

- The *dept\_count* function could instead be written as a procedure:

```
create procedure dept_count_proc (in dept_name varchar(20),
 out d_count integer)
```

```
begin
```

```
 select count(*) into d_count
 from instructor
 where instructor.dept_name = dept_count_proc.dept_name
```

```
end
```

- in** and **out** are parameters

- Values are assigned to **in** parameters by the caller
- out** parameters have values are set in the procedure in order to return results.

- Procedures can be invoked either from an SQL procedure or from embedded SQL, using the **call** statement.

```
declare d_count integer;
call dept_count_proc(‘Physics’ , d_count);
```



# SQL Procedures (Cont.)

- Procedures and functions can be invoked also from dynamic SQL
- SQL allows more than one procedure of the same name so long as the number of arguments of the procedures with the same name is different.
- The name, along with the number of arguments, is used to identify the procedure.



# Language Constructs for Procedures & Functions

- SQL supports constructs that gives it almost all the power of a general-purpose programming language (Persistent Storage Module of SQL standard)
  - Warning: most database systems implement their own variant of the standard syntax below.
- **declare** – to declare variables
- **set** – to assign values to variables
- Compound statement: **begin ... end**,
  - May contain multiple SQL statements between **begin** and **end**.
  - Local variables can be declared within compound statements
  - **begin atomic ... end** for atomic transactions



## Language Constructs for Procedures & Functions cont.

- While and repeat statements:

- **while** boolean expression **do**  
sequence of statements ;  
**end while**
  - **repeat**  
sequence of statements ;  
**until** boolean expression  
**end repeat**



# Language Constructs (Cont.)

- **for** loop
  - Permits iteration over all results of a query
- Example: Find the budget of all departments

```
declare n integer default 0;
for r as
 select budget from department /*Fetches one row at a time*/
 where dept_name = 'Music'
do
 set n = n + r.budget
end for
```

- **leave** to exit the loop //Like a break statement in C
- **iterate** starts on the next tuple, from the beginning of the loop, skipping the remaining statements



# Language Constructs – if-then-else

- Conditional statements (**if-then-else**)

```
if boolean expression
 then statement or compound statement
elseif boolean expression
 then statement or compound statement
else statement or compound statement
end if
```



# Example 1: procedure

```
create procedure GetCustomerLevel(
 in p_customerNumber int(11),
 out p_customerLevel varchar(10))
begin
 declare creditlim integer;

 select creditlimit into creditlim
 from customers
 where customerNumber = p_customerNumber;

 if creditlim > 50000 then
 set p_customerLevel = 'PLATINUM';
 elseif (creditlim <= 50000 AND creditlim >= 10000) then
 set p_customerLevel = 'GOLD';
 elseif creditlim < 10000 then
 set p_customerLevel = 'SILVER';
 end if;

end;
```



## Example 2: procedure

- Registers student for a course after ensuring classroom capacity is not exceeded
  - Returns 0 on success and -1 if capacity is exceeded
  - *takes (ID, course id, sec id, semester, year, grade)* –courses that students take
  - *classroom (building, room number, capacity)*
  - *section (course id, sec id, semester, year, building, room number, time slot id)*

-- Registers a student after ensuring classroom capacity is not exceeded  
-- Returns 0 on success, and -1 if capacity is exceeded.

```
create function registerStudent(
 in s_id varchar(5),
 in s_courseid varchar(8),
 in s_secid varchar(8),
 in s_semester varchar(6),
 in s_year numeric(4,0),
 out errorMsg varchar(100)
```

returns integer

begin

```
declare currEnrol int;
```

```
select count(*) into currEnrol
 from takes
 where course_id = s_courseid and sec_id = s_secid
 and semester = s_semester and year = s_year;
```

```
declare limit int;
```

```
select capacity into limit
 from classroom natural join section
 where course_id = s_courseid and sec_id = s_secid
 and semester = s_semester and year = s_year;
```

```
if (currEnrol < limit)
```

begin

```
 insert into takes values
```

```
 (s_id, s_courseid, s_secid, s_semester, s_year, null);
```

```
 return(0);
```

end

-- Otherwise, section capacity limit already reached

```
set errorMsg = 'Enrollment limit reached for course ' || s_courseid
 || ' section ' || s_secid;
```

```
return(-1);
```

end;



# Exception handling

- Declaring an exception:

```
declare table_not_found condition for 1051; /*my_sql_errorcode
for unknown table */
```

```
declare exit handler for table_not_found
```

```
begin
```

```
select 'Please create table abc first';
```

```
end
```

- Causing an exception:

```
create procedure print()
```

```
begin
```

```
select * from abc; //abc does not exist
```

```
end
```



# Exceptions

- Signaling of exception conditions, and declaring handlers for exceptions

```
declare <condition_name> condition for <condition_value>
declare <where_to_go> handler for <condition_name>
begin
 ...
end
```

- <condition\_value> may be
  - A well known integer value, depending on the RDBMS chosen
  - SQLSTATE <five\_character\_string>
- <where\_to\_go> may be
  - exit
  - continue



# Exceptions cont.

- The handler says that if the condition in condition\_value arises the action to be taken is to exit (continue) the enclosing **begin end** statement.
- Can raise an exception explicitly by **signal <condition\_name>**
- Pre-defined conditions
  - **sqlexception**
  - **sqlwarning**
  - **not found**

**declare continue handler for not found**

**begin**

--body of handler

**end;**



# Calling functions/procedures from JDBC

- Calling functions and procedures
  - Prepare the callable statement
    - ▶ `CallableStatement cStmt1 = conn.prepareCall("{? = call <function_name>(?)}");`
    - ▶ `CallableStatement cStmt2 = conn.prepareCall("{call <procedure_name>(?,?)");`
  - Register data types of function return values and out parameters using `registerOutParameter()`
  - Set input values (as before)
  - Execute callable statements
    - ▶ `cStmt1.execute();`
  - Retrieve results (as before)



# External Language Routines



# External Language Routines

- SQL allows us to define functions in a programming language such as Java, C#, C or C++.
  - Can be more efficient than functions defined in SQL
  - computations that cannot be carried out in SQL can be executed by these functions.
- Declaring external language procedures and functions

```
create procedure dept_count_proc(in dept_name varchar(20),
 out count integer)
```

**language** C

**external name** '/usr/avi/bin/dept\_count\_proc'

```
create function dept_count(dept_name varchar(20))
```

**returns** integer

**language** C

**external name** '/usr/avi/bin/dept\_count'



# External Language Routines (Cont.)

- Benefits of external language functions/procedures:
  - more efficient for many operations, and more expressive power.
- Drawbacks
  - Code to implement function may need to be loaded into database system and executed in the database system's address space.
    - ▶ risk of accidental corruption of database structures
    - ▶ security risk, allowing users access to unauthorized data
  - There are alternatives, which give good security at the cost of potentially worse performance.
  - Direct execution in the space of the database system is used when efficiency is more important than security.



# Security with External Language Routines

- To deal with security problems, we can do one of the following:
  - Use **sandbox** techniques
    - ▶ That is, use a safe language like Java, which cannot be used to access/damage other parts of the database code.
  - Run external language functions/procedures in a separate process, with no access to the memory of the database process.
    - ▶ Parameters and results communicated via inter-process communication
- Both have performance overheads
- Many database systems support both above approaches as well as direct executing in database system address space.



# Triggers



# Revision: Constraints discussed so far

- Domain constraints
- **not null**, **unique** and **check** constraints, **primary keys** - on attributes, on *single* relations
- Referential integrity constraints – **foreign keys**, **check** constraints across relations



# Revision: Constraints discussed so far

## ■ assertions

- A predicate that the database must always satisfy
  - ▶ Example: For each tuple in student, *tot\_cred* must equal the total credits of courses passed by the student
- Most databases do not support checks or assertions
- Hard to implement because it is unclear when to check them
  - ▶ Example: MIN must be greater (MAX must be less) than the sum of values of an attribute
  - ▶ CREATE ASSERTION SumSalary CHECK (10000 >= ALL (SELECT SUM(salary) FROM instructor GROUP BY department))
- Equivalent functionality can be implemented using triggers
- Triggers tell when the action associated with them must be executed



# Triggers

- A **trigger** is a statement that is
  - executed automatically
  - as a side effect of a modification to the database.
- More general and flexible compared to other integrity constraints



# Need for triggers

- Used to
  - Implement integrity constraints that cannot be specified using the constraint mechanism, in flexible ways
    - ▶ Preparing an order
      - A sales clerk enters name of item and ordered quantity
      - Current price is automatically inserted from another table (prevents manual errors, maintains a complete record)
  - For alerting / starting tasks automatically when certain conditions are met
    - ▶ When an inventory level falls below minimum, place an order automatically
- Can gather statistics on database modifications
  - If a user has made enough purchases in the previous month, give a discount
  - Use a separate application to alert the user



# Designing triggers

- To design a trigger mechanism, we must:
  - Specify the conditions under which the trigger is to be executed:
    - ▶ When must the trigger be checked ? Upon an *event* [this is always specified]
    - ▶ What *condition* must be satisfied ? [not just prevention]
  - Specify the *actions* to be taken when the trigger executes [prevent the event, undo its effects, perform any action, even across unrelated tables]
- Triggers introduced in SQL standard in SQL:1999
  - Syntax illustrated here may not work exactly on your database system; check the system manuals



# Triggering Events and Actions in SQL

- Triggering event can be **insert**, **delete** or **update**
- Triggers on update can be restricted to specific attributes
  - For example, **after update of takes on grade**
- Values of attributes before and after an update can be referenced
  - **referencing old row as** : for deletes and updates
  - **referencing new row as** : for inserts and updates



# Triggering Events and Actions in SQL cont.

- Triggers can be activated *before* an event, which can serve as an extra constraint (preventing invalid updates). For example, convert blank grades to null.

```
create trigger setnull_trigger before update of takes
referencing new row as nrow
for each row
 when (nrow.grade = ' ')
begin atomic
 set nrow.grade = null;
end;
```

- Event ? Condition ? Action ?



# Trigger to Maintain credits\_earned value

- create trigger *credits\_earned* after update of *takes* on (*grade*) referencing new row as *nrow* referencing old row as *orow* for each row
  - when *nrow.grade* <> 'F' and *nrow.grade* is not null and (*orow.grade* = 'F' or *orow.grade* is null)
  - begin atomic
    - update *student*
    - set *tot\_cred*= *tot\_cred* +
      - (select *credits*
      - from *course*
      - where *course.course\_id*= *nrow.course\_id*)
    - where *student.id* = *nrow.id*;
  - end;
- Event ? Condition ? Action ?



# When Not To Use Triggers

- Triggers were used earlier for tasks such as
  - Maintaining summary data (e.g., total salary of each department)
  - Replicating databases by recording changes to special relations (called **change** or **delta** relations) and having a separate process that applies the changes over to a replica
- There are better ways of doing these now:
  - Databases today provide built in materialized view facilities to maintain summary data
  - Databases provide built-in support for replication
- Encapsulation facilities can be used instead of triggers in many cases
  - Define methods to update fields
  - Carry out actions as part of the update methods instead of through a trigger



# When Not To Use Triggers (Cont.)

- Risk of unintended execution of triggers, for example, when
  - Loading data from a backup copy
  - Replicating updates at a remote site
    - ▶ Triggers are already executed on the master copy, no need to replicate
    - ▶ Trigger execution can be disabled before such actions.
- Other risks with triggers:
  - Trigger error leading to failure of critical transactions that set off the trigger
  - Cascading execution – an insert on the same relation, with a trigger on insert



# End of Module 7

Edited by Radhika Sukapuram

Database System Concepts, 7<sup>th</sup> Ed.

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use